


# LET'S FACE IT

## EMOTIONAL RECOGNITION MODEL

Georgiy Sekretaryuk, Rafael Arbex-Murut, Yeshwanth Somu



# PROJECT OVERVIEW

## THE DATA

- This [Kaggle Facial Recognition Dataset](#) contains grayscale images with different facial expressions.
  - surprise, anger, happiness, sad, neutral, fear
- Training Set: 28,079 (80%) // Testing set: 7,178 (20%)
- Attributes: 2304 (48 x 48 pixels)

## OUR PROBLEM

Our problem was - how do we build a good enough model to predict emotions (that may vary in perception)?

We started with an attempt to build an FNN model first and assess performance. We decided that if performance was low, we would build a Convolutional Neural Network Model (CNN).




*How do we predict emotions?*



# INITIAL BASELINE MODEL - FNN



- The model performed better than random guessing would



# DIFFERENT STRATEGY - CNN MODEL

- We decided to use a CNN Model instead. As we learned in the course, CNN models account for inconsistencies in images. While our images were of the same size, the faces were certainly not standardized like the MNIST dataset.
- Different zooms, face sizes, etc.

Layer (type)	Output Shape	Param #
conv_1 (Conv2D)	(None, 48, 48, 32)	832
pool_1 (MaxPooling2D)	(None, 24, 24, 32)	0
conv_2 (Conv2D)	(None, 24, 24, 64)	51264
pool_2 (MaxPooling2D)	(None, 12, 12, 64)	0
flatten_1 (Flatten)	(None, 9216)	0
fc_1 (Dense)	(None, 1024)	9438208
dropout_1 (Dropout)	(None, 1024)	0
fc_2 (Dense)	(None, 6)	6150
Total params: 9,496,454		
Trainable params: 9,496,454		
Non-trainable params: 0		



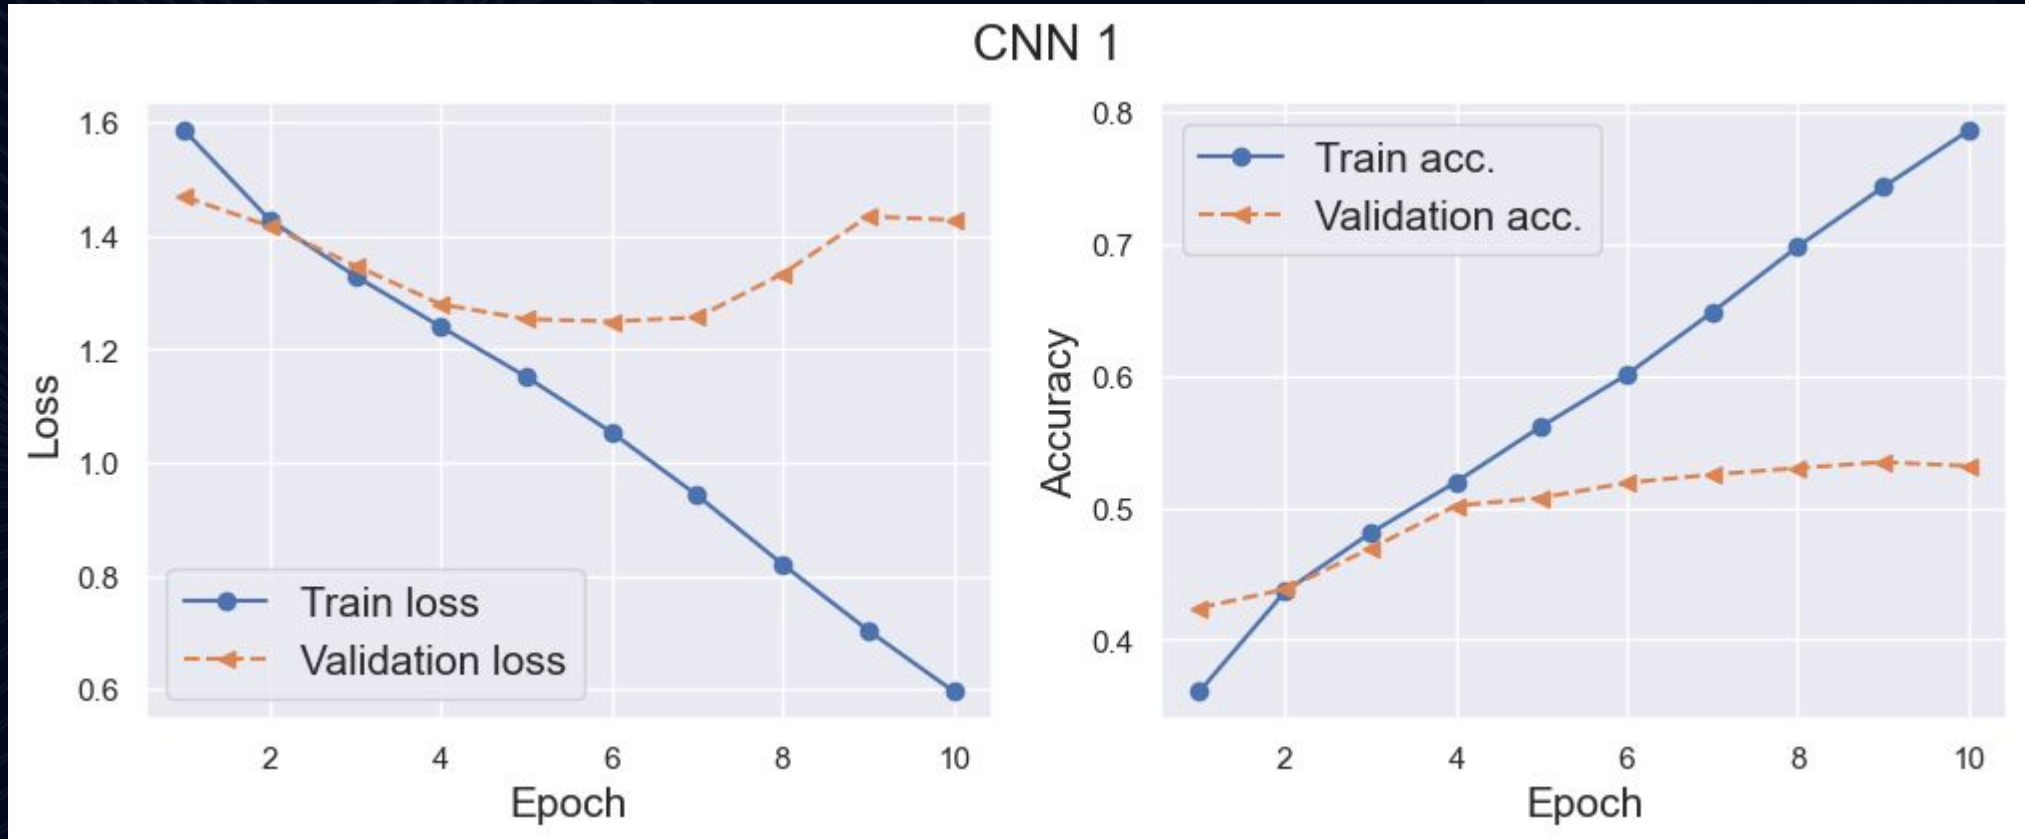
# DIFFERENT STRATEGY - CNN MODEL

...RESULTS IMPROVED

Epoch	Loss	Accuracy	Val Loss	Val Accuracy
1	1.5860	0.3612	1.4623	0.4208
2	1.4091	0.4413	1.3460	0.4694
3	1.2945	0.4984	1.2824	0.4974
4	1.2018	0.5389	1.2534	0.5131
5	1.1008	0.5828	1.2163	0.5353
6	0.9949	0.6255	1.2260	0.5373
7	0.8701	0.6776	1.2473	0.5397
8	0.7450	0.7249	1.3136	0.5455
9	0.6271	0.7702	1.3630	0.5541
10	0.5250	0.8119	1.4674	0.5538



# DIFFERENT STRATEGY - CNN MODEL



True Label \ Predicted Label	Happy	Surprise	Angry	Neutral	Fear	Sad
Happy	1338	41	47	162	76	110
Surprise	46	605	22	58	73	27
Angry	108	28	353	185	108	176
Neutral	131	26	55	716	95	210
Fear	88	90	96	168	374	208
Sad	145	29	110	282	153	528

	precision	recall	f1-score	support
0	0.721	0.754	0.737	1774
1	0.739	0.728	0.733	831
2	0.517	0.368	0.430	958
3	0.456	0.581	0.511	1233
4	0.425	0.365	0.393	1024
5	0.419	0.423	0.421	1247
accuracy			0.554	7067
macro avg	0.546	0.537	0.538	7067
weighted avg	0.553	0.554	0.550	7067

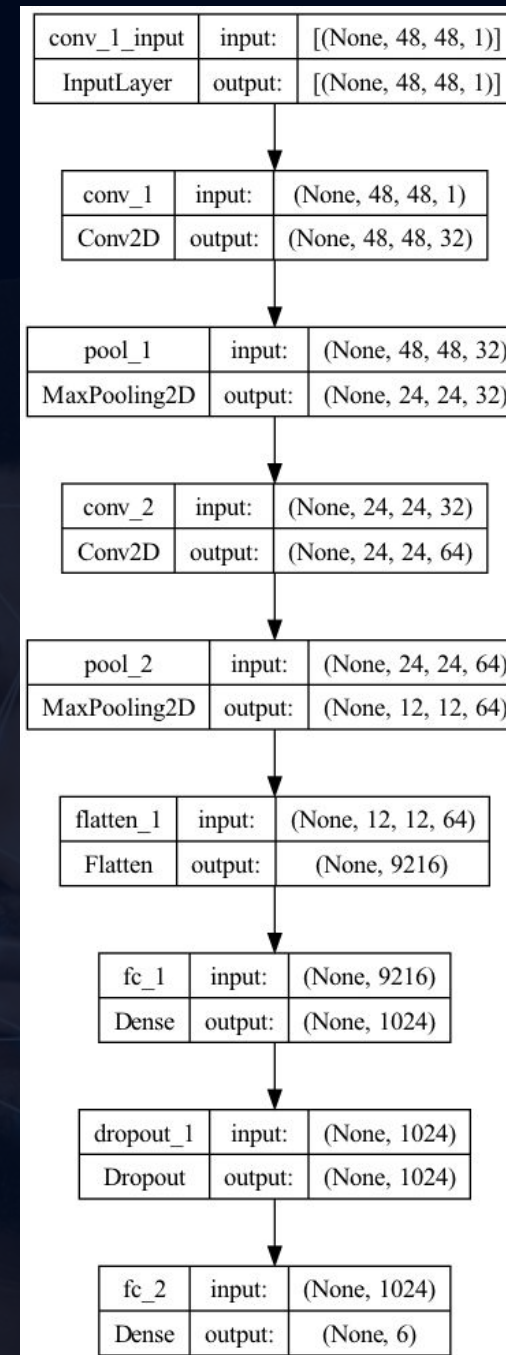
# 2ND MODEL - TWEAKING MODEL 1

- We decided to adjust Model 1 and try:
  - a different optimizer (SGD)
  - Adjusting the kernel size to 4x4 instead of 5x5 pixels.

Layer (type)	Output Shape	Param #
conv_1 (Conv2D)	(None, 48, 48, 32)	544
pool_1 (MaxPooling2D)	(None, 24, 24, 32)	0
conv_2 (Conv2D)	(None, 24, 24, 64)	32832
pool_2 (MaxPooling2D)	(None, 12, 12, 64)	0
flatten_1 (Flatten)	(None, 9216)	0
fc_1 (Dense)	(None, 1024)	9438208
dropout_1 (Dropout)	(None, 1024)	0
fc_2 (Dense)	(None, 6)	6150

---

Total params: 9477734 (36.15 MB)
Trainable params: 9477734 (36.15 MB)
Non-trainable params: 0 (0.00 Byte)





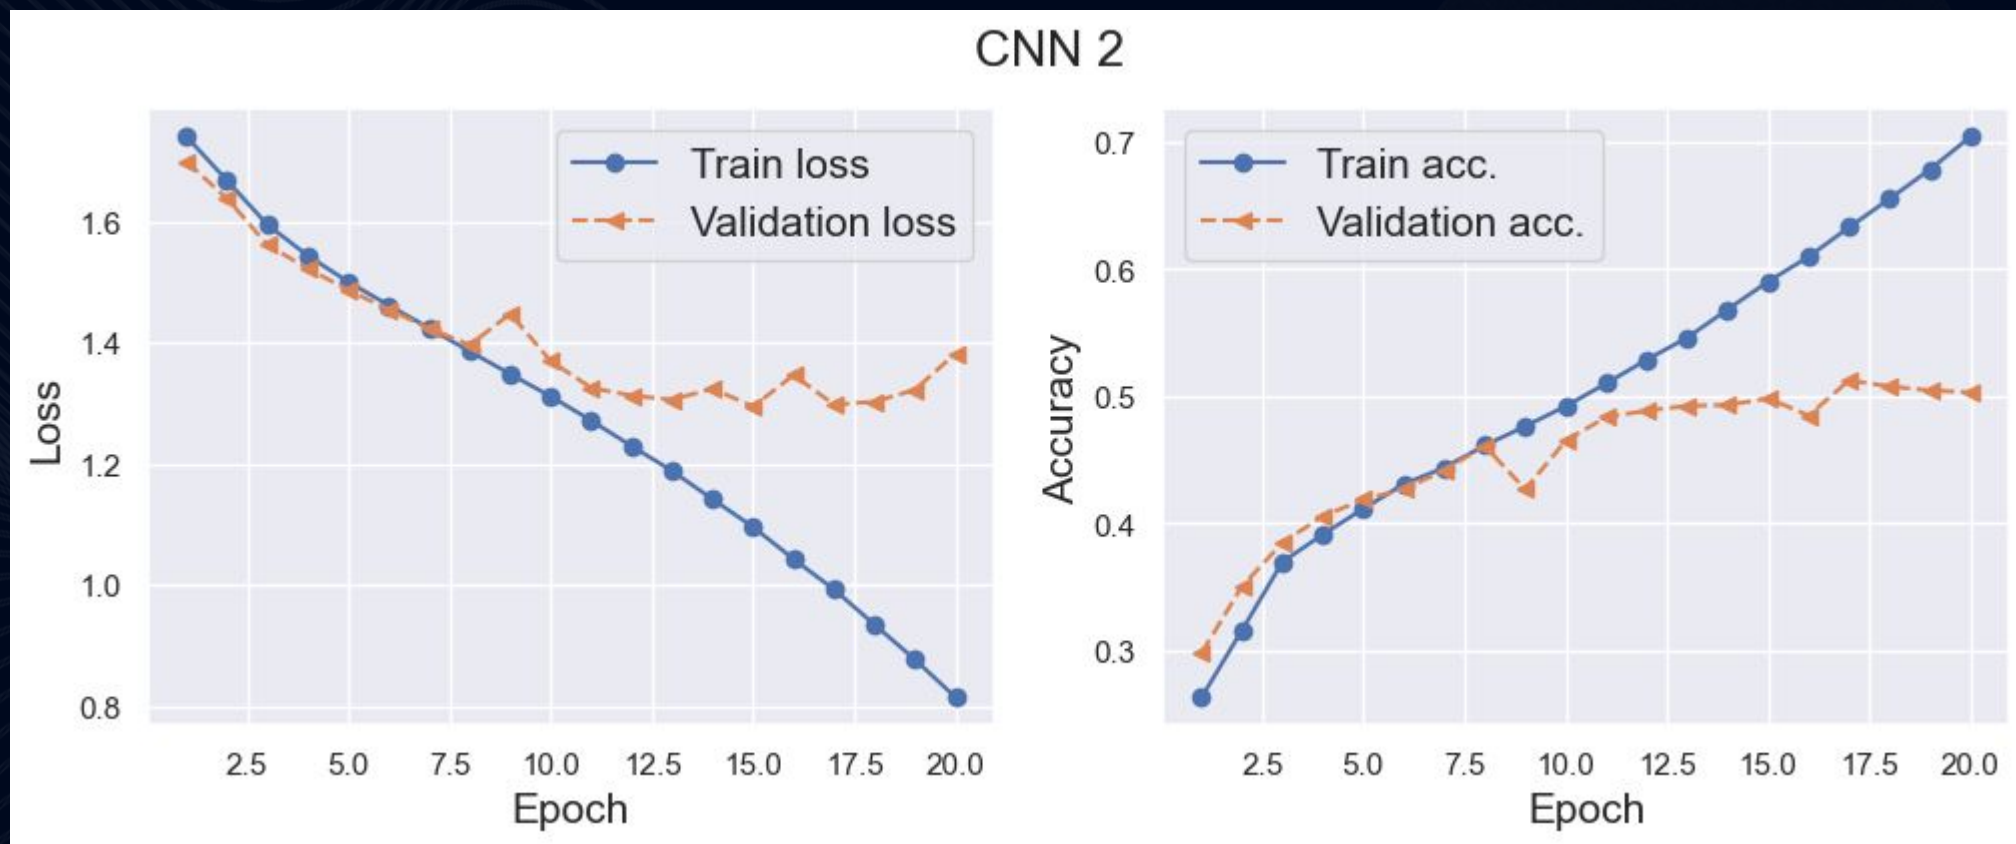
# AND THE RESULTS FOR MODEL 2...

DIDN'T IMPROVE.

Epoch	Loss	Accuracy	Val Loss	Val Accuracy
1	1.5584	0.3712	1.4057	0.4439
2	1.3557	0.4685	1.3015	0.4900
3	1.2409	0.5204	1.2580	0.5110
4	1.1337	0.5691	1.2330	0.5299
5	1.0233	0.6151	1.2335	0.5370
6	0.8933	0.6690	1.2322	0.5489
7	0.7547	0.7249	1.2941	0.5427
8	0.6264	0.7709	1.3548	0.5544
9	0.5131	0.8149	1.4346	0.5536
10	0.4264	0.8465	1.5872	0.5531

# 2ND MODEL PERFORMANCE

## VALIDATION LOSS AND ACCURACY SUFFERED



# 2ND MODEL - WORSE THAN MODEL 1

True Label	Happy	Sad	Fear	Neutral	Angry	Suprise
	1452	93	40	115	42	32
	84	685	141	203	106	28
	68	272	382	99	102	101
	128	304	91	608	66	36
	69	217	108	101	440	23
	62	18	114	19	18	600
	Happy	Sad	Fear	Neutral	Angry	Suprise
	Predicted Label					

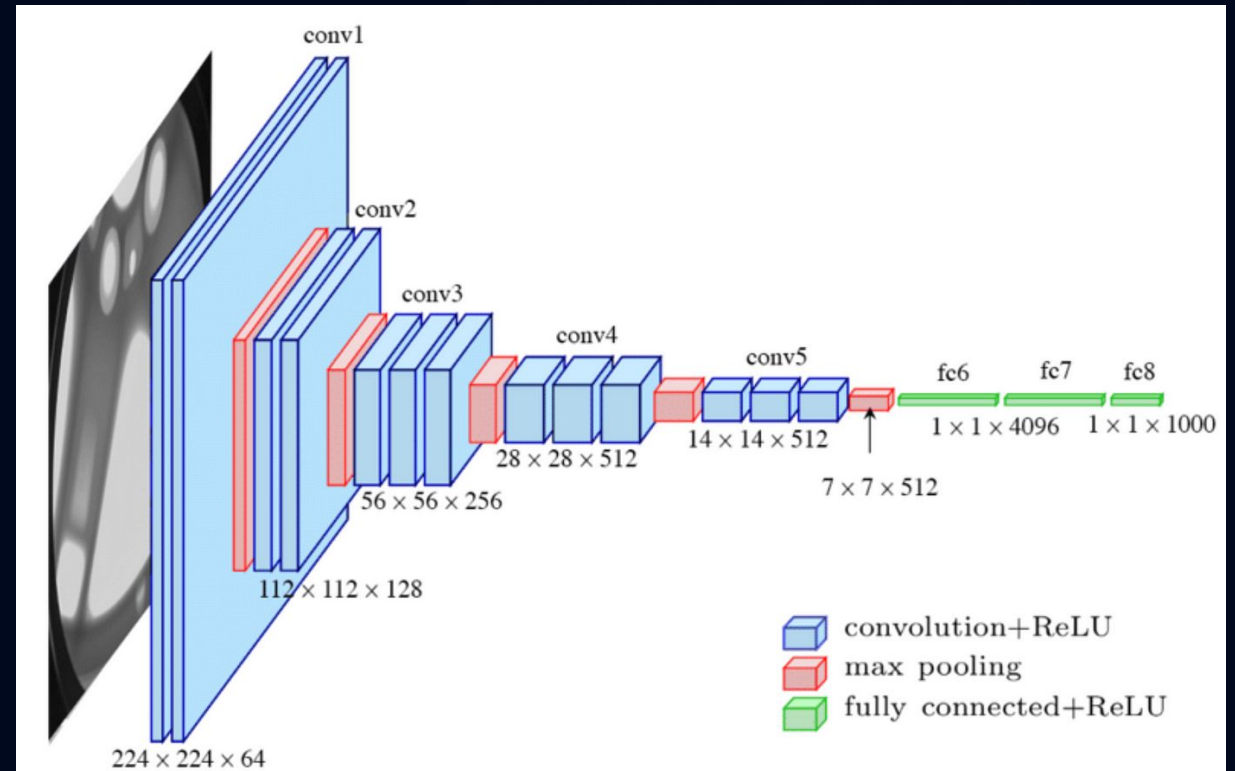
	precision	recall	f1-score	support
0	0.531	0.702	0.605	1774
1	0.600	0.005	0.010	1247
2	0.346	0.225	0.273	1024
3	0.345	0.397	0.370	1233
4	0.368	0.281	0.319	958
5	0.340	0.777	0.473	831
accuracy			0.408	7067
macro avg	0.422	0.398	0.341	7067
weighted avg	0.439	0.408	0.356	7067



# LET'S TRY A DIFFERENT ARCHITECTURE...

## ENTER VGG 16

- VGG 16 is a CNN architecture used to win the Imagenet competition in 2014.
- Considered to be a high performing vision model.
- Focuses on convolutional layers of 3x3 filters with a stride of 1, always uses the same padding, and a maxpool layer of 2x2 filter with a stride of 2.
- Has 16 layers (hence the name).
- Approximately 138 million parameters.



# VGG 16 TRAINING

Epoch	Loss	Accuracy	Val Loss	Val Accuracy
1	1.6487	0.3093	1.4372	0.431
2	1.3526	0.4551	1.2756	0.4868
3	1.1843	0.5332	1.1727	0.5373
4	1.0539	0.5909	1.1011	0.5714
5	0.9257	0.6452	1.0988	0.586
6	0.7773	0.7076	1.1524	0.593
7	0.6197	0.7707	1.1552	0.6029
8	0.4638	0.8328	1.3542	0.5981
9	0.3201	0.8871	1.5211	0.5821
10	0.231	0.9204	1.6232	0.586

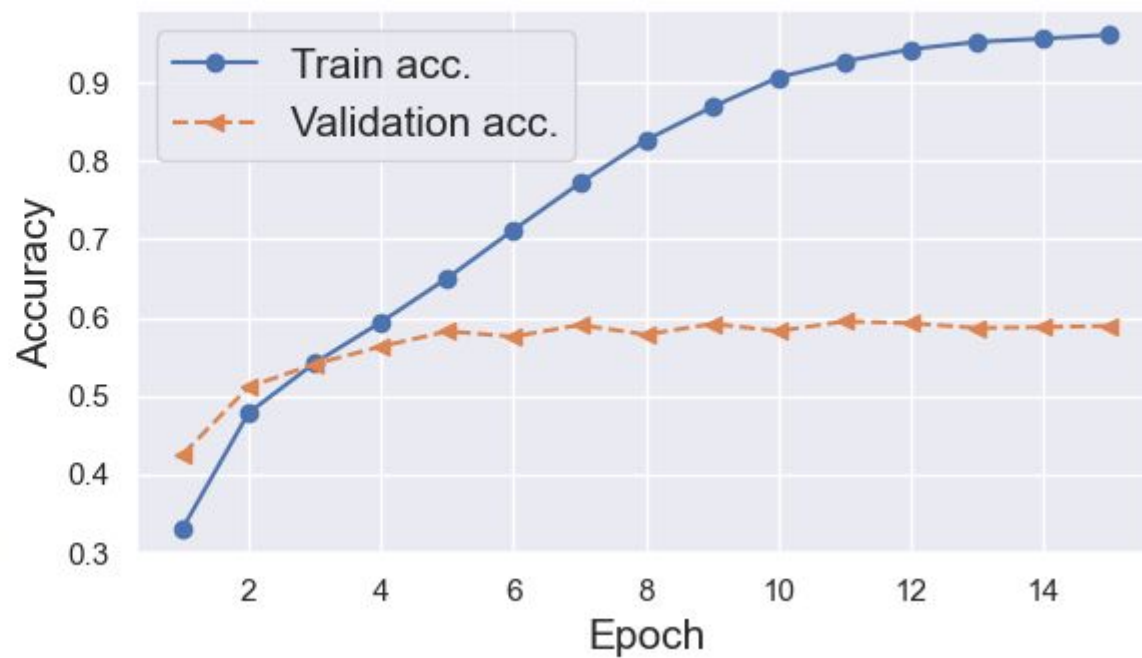

## RESULTS IMPROVED, SLIGHTLY

True Label \ Predicted Label	Happy	Sad	Fear	Neutral	Angry	Suprise
Happy	1452	93	40	115	42	32
Sad	84	685	141	203	106	28
Fear	68	272	382	99	102	101
Neutral	128	304	91	608	66	36
Angry	69	217	108	101	440	23
Suprise	62	18	114	19	18	600





# PERFORMANCE WITH VGG 16

VGG16-based Model



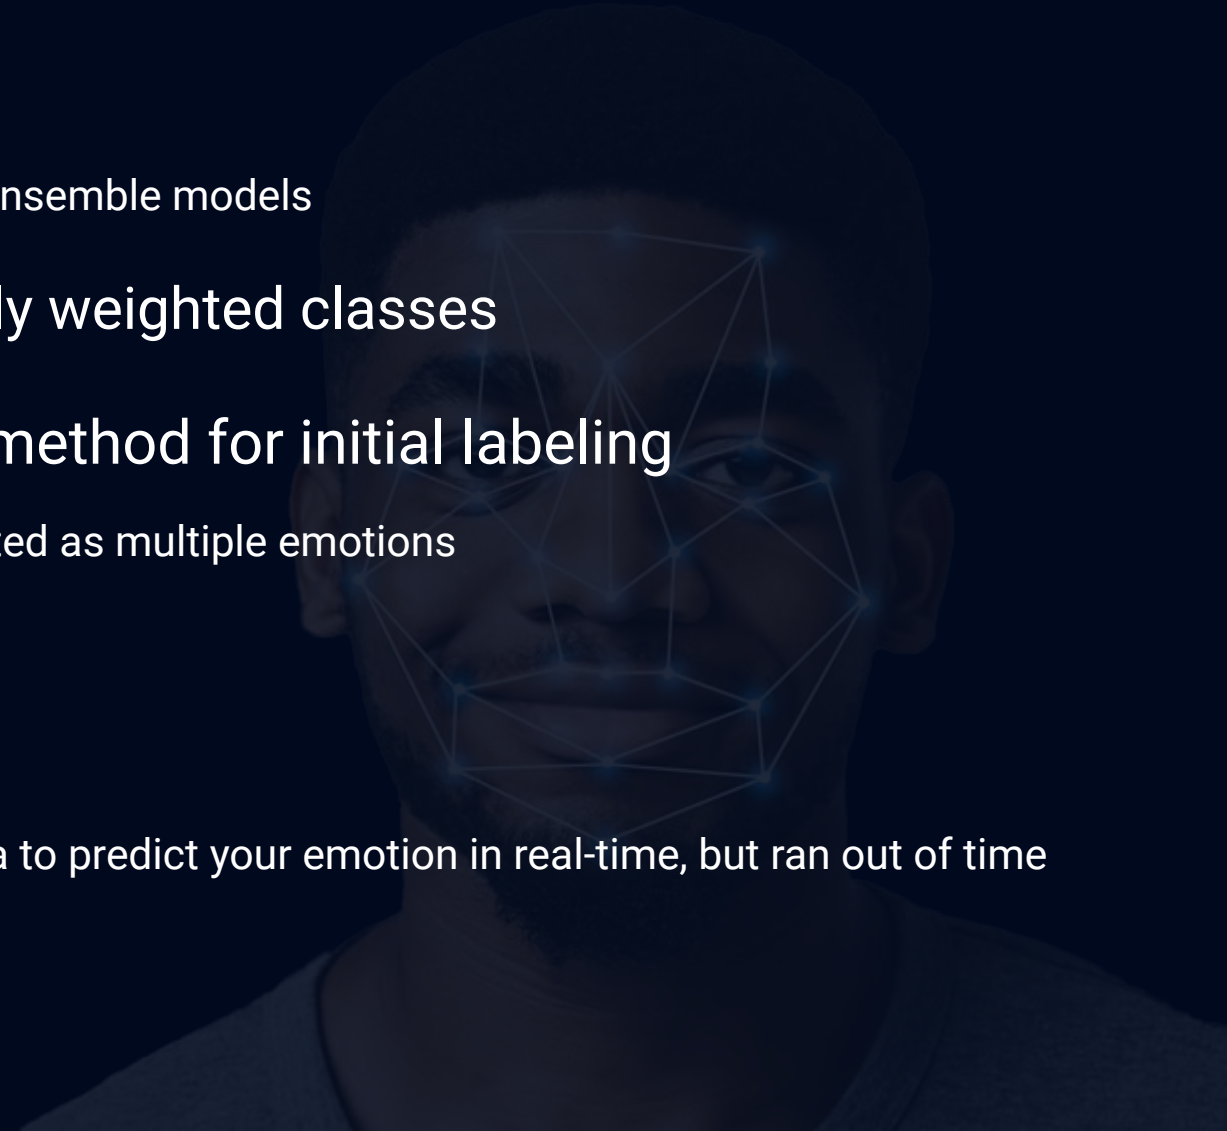
# LEARNING

- Model architecture learnings
  - Model 2 did not perform better with SGD / smaller Kernels
  - VGG 16 model performed better
- Lot of happy faces in training data
  - 25% of the training data was happy faces, leading to unequal classes
  - Overfitting on the “Happy” faces in some cases
- Location of faces in images matters
  - Non-centered faces were predicted poorly
- Emotion labeling can be subjective
  - Many faces can be interpreted as different emotions such as Neutral OR Angry



# If we had more time

- Explore other model architectures
  - Potentially more parameters, more kernel sizes, or ensemble models
- Work on larger data sets with more equally weighted classes
- Have more emotions and define better method for initial labeling
  - Especially for expressions that can be interpreted as multiple emotions
  - Additional emotions: Doubt, Pain, Disgust
- Connect to camera and turn into an app
  - Explored creating a Flask app that uses your camera to predict your emotion in real-time, but ran out of time





The background features a dark blue gradient. On the left, there are light blue wavy lines. On the right, there is a faint image of a person's face with a blue wireframe overlay, suggesting a digital or AI theme.

# THANK YOU!

**HAPPY TO ANSWER ANY ADDITIONAL QUESTIONS**